

## Chapter 3

### Arithmetic

- Digits are added bit-by-bit from right to left, with carries passed to the next digit to the left
- Subtraction is simply addition with the appropriate operand negated before the addition
- Overflow occurs when the result from an operation cannot be represented with the available hardware
  - Overflow cannot occur in addition with numbers of differing signs because the result is no larger than one of the operands
    - Therefore, in subtraction, overflow cannot occur when the signs are the same
  - Adding/subtracting two 32-bit numbers can yield a result that needs 33 bits
    - We don't have 33 bits, so that means when overflow occurs, the sign bit is set with the value of the result instead of the proper sign

### Overflow

- Therefore, overflow occurs when adding two positive numbers and the sum is negative, or vice-versa
  - In subtraction, it occurs when we subtract a negative number from a positive number and get a negative result or subtract a positive number from a negative number and get a positive result
- The above is for signed (Two's Complement) overflow. For unsigned overflow, we usually choose to ignore it, therefore MIPS includes two types of arithmetic instructions:
  - *add*, *addi*, and *sub* cause exceptions on overflow (signed)
  - *addu*, *addiu*, and *subu* do *not* cause exceptions on overflow
- C always chooses to ignore overflow so the compiler will always generate the unsigned versions of the instructions
- If we were to detect overflow, MIPS would do so with an exception/interrupt, which is an unscheduled event that disrupts program execution
  - The address of the instruction that overflowed is saved in a register and the computer jumps to a predefined address to invoke the appropriate routine for that exception
    - MIPS includes a register called exception program counter (EPC) that contains that address of the instruction that caused the exception so that the software can return to it if it wants to
      - This also leads to the existence of \$k0 and \$k1 registers reserved for the operating system so that it can be used when jumping back to the offending instruction and resetting all general-purpose registers without using a general purpose register to jump back with

## Chapter 3

Floating Points	<ul style="list-style-type: none"><li>• We also need support for fractions (real numbers) and not just signed and unsigned numbers</li><li>• A number like 3,155,760,000 is too large to be stored in 32 bits, but can be represented in a more compact form using scientific notation in the form of <math>3.15576 \times 10^9</math><ul style="list-style-type: none"><li>• A normalized number is a number in floating-point notation that has no leading 0s</li></ul></li><li>• We can also show binary numbers in scientific notation</li><li>• Floating point refers to computer arithmetic that represents numbers in which the binary point is not fixed<ul style="list-style-type: none"><li>• C uses the name <i>float</i> for such numbers</li></ul></li></ul>
Fields	<ul style="list-style-type: none"><li>• Must find a compromise between the size of the fraction and the exponent<ul style="list-style-type: none"><li>• Increasing the size of the fraction enhances the precision of the fraction</li><li>• Increasing the size of the exponent increases the range of numbers that can be represented</li></ul></li><li>• Floating point numbers are of the form: <math>(-1)^S \times F \times 2^E</math></li><li>• For a 32 bit floating point number, we have:<ul style="list-style-type: none"><li>• <i>Sign (1 bit)</i>: sign of the fraction (1 if it is negative, 0 if it is positive)</li><li>• <i>Exponent (8 bits)</i>: Value of the exponent; includes sign</li><li>• <i>Fraction (23 bits)</i>: This is a 23-bit number</li></ul></li></ul>
Range	<ul style="list-style-type: none"><li>• These sizes give an extraordinary range:<ul style="list-style-type: none"><li>• Fractions can be almost as small as <math>2.0 \times 10^{-38}</math> and numbers almost as large as <math>2.0 \times 10^{38}</math></li></ul></li><li>• Numbers can still be too large<ul style="list-style-type: none"><li>• Overflow occurs when exponent is too large to be represented in the exponent field</li></ul></li><li>• We also have issues of numbers being too small<ul style="list-style-type: none"><li>• <b>Underflow</b> occurs when the negative exponent has become too large to fit in the exponent field</li></ul></li></ul>
Double Precision	<ul style="list-style-type: none"><li>• We can reduce the number of overflow/underflow cases by introducing a double precision floating point arithmetic<ul style="list-style-type: none"><li>• It is represented in two 32-bit words (therefore it is 64 bits)</li><li>• In C, this number is called a <i>double</i></li></ul></li><li>• In double-precision, the breakdown is the following:<ul style="list-style-type: none"><li>• <i>Sign</i>: 1 bit</li><li>• <i>Exponent</i>: 11 bits</li><li>• <i>Fraction</i>: 20 bits</li><li>• This thus allows numbers as small as <math>2.0 \times 10^{-308}</math> and as large as <math>2.0 \times 10^{308}</math></li></ul></li><li>• These formats are all part of the <i>IEEE 754 floating-point standard</i> that is virtually universal to all computers</li></ul>
Normalized	<ul style="list-style-type: none"><li>• IEEE 754 makes the leading 1-bit of the normalized binary numbers implicit</li></ul>

## Chapter 3

### Special Cases

- There is leading 1-bit in the normalized version because in the normalized version of scientific notation, there is only 1 digit to the left of the decimal point. Because this is binary, the only number to the left of the binary point is a 1. Therefore, to save 1 bit we can make this implicit.
- The significand is therefore 24 bits for single-precision and 53 bits for double-precision
  - Significand refers to these lengths whereas fraction refers to the 23-bit and 52-bit numbers
- 0 has no implicit 1, so to represent this, the fraction must be all 0 and the exponent must be all 0
- For the rest of the numbers, the modified representation includes the hidden 1:  $(-1)^S \times (1 + \text{Fraction}) \times 2^E$ 
  - Numbering the bits of the fraction from left to right  $s_1, s_2, s_3, \dots$ , then the value is:
    - $(-1)^S \times (1 + (s_1 \times 2^{-1}) + (s_2 \times 2^{-2}) + (s_3 \times 2^{-3}) + \dots) \times 2^E$
- An exponent of 255 (the largest exponent) and a fraction of 0 yields either positive or negative infinity
  - Programmers can use this in instances such as dividing by 0 instead of interrupting or throwing an error
- An exponent of 255 (the largest exponent) and a nonzero fraction produces a NaN (Not a Number)
  - This is used as a result of invalid operations, such as  $0/0$  or subtracting infinity from infinity
- An exponent of 0 and a nonzero fraction is a de-normalized number
  - In this case, there is no implicit 1, which goes hand in hand with the definition of a de-normalized number
    - Therefore, we are able to reach even smaller numbers, because previously the smallest number was  $(1 + 0.00\dots00) \times 2^{-126}$  because of the implicit 1
    - Once taken out, with denorms, the smallest number is then  $0.00\dots01 \times 2^{-126} = 2^{-23} \times 2^{-126}$

### Bias Encoding

- The designers wanted a floating-point representation that could easily be processed by integer comparisons, especially sorting
  - The most significant bit is the sign bit which allows for a quick test of less than, greater than, or equal to 0
  - The exponent bit comes before the significand because numbers with bigger exponents look larger than numbers with smaller exponents as long as both exponents have the same sign
    - We can't use two's complement here because negative exponents have a 1 in the most significant bit of the exponent field, which makes it look like a big number when it's actually the opposite
    - Instead, we encode the exponent with bias encoding
      - Represent the most negative exponent as all 0s
      - The most positive exponent as all 1s
        - For single precision, use a bias of 127

- The real form is then  $(-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$
- Because the exponent of 255 is reserved for NaN's and infinity, the biggest exponent we can have is 254 which yields an upper bound of 127 on the exponent and a lower bound of -126